



DOZIER
TECH GROUP

ROLE-SPECIFIC, EVIDENCE-FIRST WRITING

GRANT DOZIER

Cover Letter & Outreach System

Reusable letters that sound like an engineer—not a template generator.

R U L E S

- Open with the specific system or mission connection, not excitement about the company.
- Use two or three pieces of evidence; never summarize the entire résumé.
- Name the transfer: explain why past architecture predicts performance in this role.
- Keep the core letter between 250 and 400 words.
- Rewrite the first paragraph and evidence selection for every application.

E V I D E N C E B A N K

- Agent infrastructure: typed MCP tools across engineering, communications, files, knowledge, transcription, and business operations.
- Applied AI at scale: ~22.5M POS units with deterministic resolution, embeddings, verification, confidence gates, and human review.
- Governed AI: ATHENA/HERMES provenance, bi-temporal rules, visible uncertainty, counter-evidence, and human-final action.
- Operational outcomes: 95% cycle-time reduction, 300 files per run, 10–20 hours recovered per person per project.



DOZIER
TECH GROUP

OPEN AI · APPLIED AI / AGENT PLATFORM

GRANT DOZIER

Cover Letter

For roles where models meet tools, infrastructure, products, and users.

Hiring Team,

The work I am strongest at sits between frontier models and real operations: the tools, retrieval, data contracts, evaluation gates, and human workflows that determine whether an AI system becomes dependable infrastructure or remains an impressive demo.

I built a typed MCP operating layer that connects engineering and business workflows across GitHub, Linear, Microsoft 365, Google Drive, Cloudflare R2, transcription, billing, and institutional memory. I also lead an AI-assisted data platform resolving approximately 22.5 million POS units through deterministic evidence, embeddings, model verification, confidence thresholds, and human review. In both systems, the central problem is not simply calling a model—it is designing the surrounding architecture so uncertainty is visible, actions are governed, and operators can trust the result.

That approach aligns with engineering work at OpenAI where well-designed APIs and abstractions must help product teams move quickly without sacrificing reliability. My background is unusually broad: founder-level product ownership, enterprise and government delivery, cloud migration, and hands-on engineering across Python, TypeScript, C#, React, SQL, .NET, and FastAPI. I am comfortable entering an ambiguous domain, finding the decision that matters, and owning the system through production.

I would welcome the opportunity to discuss how that combination of applied-AI judgment, platform engineering, and customer consequence could contribute to the team.

Grant Dozier



DOZIER
TECH GROUP

ANTHROPIC · MEMBER OF TECHNICAL STAFF

GRANT DOZIER

Cover Letter

For product, platform, safeguards, or enterprise-deployment engineering.

Hiring Team,

Anthropic's emphasis on direct evidence of ability resonates with how I have built my career: by taking consequential, poorly structured problems and turning them into working systems.

My most relevant independent work is ATHENA and HERMES, an institutional-intelligence architecture for higher education and sports organizations. The core design is a sourced decision surface: every conclusion retains provenance, applicable rule versions, counter-evidence, uncertainty, and a path for human challenge. The system is architecturally prevented from silently converting model output into an irreversible action. That work reflects the kind of engineering judgment I want to bring to increasingly capable AI systems—useful, legible, empirically grounded, and appropriately constrained.

I have also built a typed MCP operating layer spanning more than 100 TypeScript modules and lead an AI-assisted entity-resolution platform operating across approximately 22.5 million POS units. These projects required retrieval, tool design, data modeling, confidence gates, human review, and production ownership across multiple languages and services.

I am an engineer rather than a pretraining researcher. My value is in making powerful models useful and trustworthy inside complex products and organizations. I would be excited to contribute that systems perspective, high agency, and founder-level ownership to Anthropic.

Grant Dozier



DOZIER
TECH GROUP

RECRUITER · REFERRAL · HIRING MANAGER

GRANT DOZIER

Outreach Templates

Short messages designed to earn the next conversation.

RECRUITER INTRO

I'm a founder-engineer with 9+ years building applied AI and platform systems. Recent work includes a typed MCP operating layer, human-governed entity resolution across ~22.5M POS units, and institutional-intelligence architectures built around provenance and visible uncertainty. I'm targeting Staff applied-AI/platform roles and hands-on engineering leadership. Portfolio: [link].

REFERRAL REQUEST

I'm considering [role] because the team's work on [specific problem] overlaps directly with [one project]. My strongest evidence is [metric/system], and I've included a two-page résumé plus a short technical portfolio. If the fit looks real to you, would you be comfortable referring me? No pressure if the level or scope does not line up.

HIRING MANAGER NOTE

Your team's work on [specific system] caught my attention. I've built adjacent infrastructure around [tools/retrieval/evals/data/human review], including [specific evidence]. The technical portfolio below maps the claims to architecture and repositories. I'd value a conversation if the role needs someone who can own the boundary between models and production operations.